

```

1  """
2  EECS 445 - Introduction to Machine Learning
3  HW1 Q5 Perceptron Algorithm with Offset
4  """
5
6  import numpy as np
7  from helper import load_data
8
9  def all_correct(X, y, theta, b):
10     """
11     Args:
12         X: np.array, shape (n, d)
13         y: np.array, shape (n,)
14         theta: np.array, shape (d,), normal vector of decision boundary
15         b: float, offset
16
17     Returns true if the linear classifier specified by theta and b correctly classifies all examples
18     """
19
20     for i in range(len(X)):
21         for j in range(len(X)):
22             pred = np.dot(theta, X[i]) + b
23             if y[i] * pred <= 0:
24                 return False
25
26     return True
27
28
29 def perceptron(X, y):
30     """
31     Implements the Perceptron algorithm for binary linear classification.
32     Args:
33         X: np.array, shape (n, d)
34         y: np.array, shape (n,)
35
36     Returns:
37         theta: np.array, shape (d,)
38         b: float
39         alpha: np.array, shape (n,).
40             Misclassification vector, in which the i-th element is has the number of times
41             the i-th point has been misclassified)
42     """
43
44     n, d = X.shape
45     theta = np.zeros((d,))
46     b = 0.0
47     alpha = np.zeros((n,))
48
49     updated = True
50     while updated:
51         updated = False
52         for i in range(n):
53             margin = y[i] * (np.dot(theta, X[i]) + b)
54             if margin <= 0:
55                 theta += y[i] * X[i]
56                 b += y[i]
57                 alpha[i] += 1
58                 updated = True
59
60     return theta, b, alpha
61
62
63
64 def main(fname):
65     X, y = load_data(fname)
66     theta, b, alpha = perceptron(X, y)
67
68     print("Done!")
69     print("==== Classifier =====")
70     print("Theta: ", theta)
71     print("b: ", b)
72
73     print("\n")
74     print("==== Alpha =====")
75     print("i \t Number of Misclassifications")
76     print("=====")
77     for i in range(len(alpha)):
78         print(i, "\t\t", alpha[i])
79     print("Total Number of Misclassifications: ", np.sum(alpha))
80
81
82 if __name__ == '__main__':
83     main('dataset/q5.csv')

```

Q5

```

8 import numpy as np
9 from helper import load_data
10 import time
11
12 def calculate_squared_loss(X, y, theta):
13     """
14     Args:
15         X: np.array, shape (n, d)
16         y: np.array, shape (n,)
17         theta: np.array, shape (d,). Specifies an (d-1)^th degree polynomial
18
19     Returns:
20         The squared loss for the given data and parameters
21     """
22     n, _ = X.shape
23     r = X @ theta - y
24     return 0.5 * np.dot(r, r) / n
25
26 def ls_gradient_descent(X, y, learning_rate=0):
27     """
28     Implement the Gradient Descent (GD) algorithm for least squares regression.
29     Note:
30         - Please use the following stopping criteria together: number of iterations >= 1e6 or |new_loss - prev_loss| <= 1e-10
31
32     Args:
33         X: np.array, shape (n, d)
34         y: np.array, shape (n,)
35
36     Returns:
37         theta: np.array, shape (d,)
38     """
39     _, d = X.shape
40     theta = np.zeros(d)
41
42     eps = 1e-10
43     max_iter = 1e6
44     n_iter = 0
45
46     step = learning_rate
47     prev_loss = np.inf
48     new_loss = calculate_squared_loss(X, y, theta)
49
50
51
52     while (n_iter < max_iter) and (abs(new_loss - prev_loss) > eps):
53         n_iter += 1
54         grad = []
55
56         for xt, yt in zip(X, y):
57             pred = np.dot(theta, xt)
58             g = (pred - yt) * xt
59             grad.append(g)
60
61         gradient = np.mean(grad, axis=0)
62         theta -= gradient * step
63
64         prev_loss = new_loss
65         new_loss = calculate_squared_loss(X, y, theta)
66
67     print("Learning rate:", learning_rate, "\t\t\tNum iterations:", n_iter)
68     return theta
69

```

Q6

```

73 def ls_stochastic_gradient_descent(X, y, learning_rate=0):
74     """
75     Implement the Stochastic Gradient Descent (SGD) algorithm for least squares regression.
76     Note:
77     | - Please do not shuffle your data points.
78     | - Please use the following stopping criteria together: number of iterations >= 1e6 or |new_loss - prev_loss| <= 1e-10
79
80     Args:
81     | X: np.array, shape (n, d)
82     | y: np.array, shape (n,)
83     | learning_rate: the learning rate for the algorithm
84
85     Returns:
86     | theta: np.array, shape (d,)
87     """
88     _, d = X.shape
89     theta = np.zeros(d)
90     adaptive = (learning_rate == 'adaptive')
91
92     eps = 1e-10
93     max_iter = 1e6
94     n_iter = 0
95     epochs = 0
96
97     step = learning_rate
98     prev_loss = np.inf
99     new_loss = calculate_squared_loss(X, y, theta)
100
101     if adaptive:
102         step = 0.25 # best results
103     else:
104         step = 0.01 if learning_rate == 0 else float(learning_rate)
105
106     while (n_iter < max_iter) and (abs(new_loss - prev_loss) > eps):
107         epochs += 1
108
109         # simply decrementing our learning rate as we go...gradually approaches 0
110         if adaptive:
111             step /= (epochs + 1)
112
113         for xt, yt in zip(X, y):
114             pred = np.dot(theta, xt)
115             gradient = (pred - yt) * xt
116
117             theta -= step * gradient
118             n_iter += 1
119
120             if n_iter >= max_iter:
121                 break
122
123         prev_loss = new_loss
124         new_loss = calculate_squared_loss(X, y, theta)
125
126     print("Learning rate:", learning_rate, "\t\t\tNum iterations:", n_iter)
127     return theta
128
129
130
131
132 def ls_closed_form_optimization(X, y):
133     """
134     Implement the closed form solution for least squares regression.
135
136     Args:
137     | X: np.array, shape (n, d)
138     | y: np.array, shape (n,)
139
140     Returns:
141     | theta: np.array, shape (d,)
142     """
143     # piazza @85
144     return np.linalg.pinv(X) @ y

```

Q6

Q6

```
def main(fname_train):
    # TODO: This function should contain all the code you implement to complete question 6.

    X_train, y_train = load_data(fname_train)

    # Appending a column of constant ones to the X_train matrix to make X_train the same dimensions as theta.
    # The term multiplied by theta_0 is  $x^0 = 1$  (theta_0 is a constant), which is why the column contains only ones.
    X_train = np.hstack((np.ones((X_train.shape[0], 1)), X_train))

    print("Closed-form:")
    t11 = time.process_time()
    theta_cf = ls_closed_form_optimization(X_train, y_train)
    t1 = t11 - time.process_time()
    loss_cf = calculate_squared_loss(X_train, y_train, theta_cf)
    print("Squared Loss:", loss_cf)
    print("Theta", np.linalg.norm(theta_cf))
    print("\n")

    print("Gradient Decent with LR=0.01:")
    t21 = time.process_time()
    theta_gd = ls_gradient_descent(X_train, y_train, learning_rate=0.01)
    t2 = t21 - time.process_time()
    loss_gd = calculate_squared_loss(X_train, y_train, theta_gd)
    print("Squared Loss:", loss_gd)
    print("Theta", np.linalg.norm(theta_gd))
    print("\n")

    print("Gradient Decent with LR=0.05:")
    t31 = time.process_time()
    theta_gd2 = ls_gradient_descent(X_train, y_train, learning_rate=0.05)
    t3 = t31 - time.process_time()
    loss_gd2 = calculate_squared_loss(X_train, y_train, theta_gd2)
    print("Squared Loss:", loss_gd2)
    print("Theta", np.linalg.norm(theta_gd2))
    print("\n")

    print("Stochastic Gradient Decent with LR=0.01:")
    t41 = time.process_time()
    theta_sgd = ls_stochastic_gradient_descent(X_train, y_train, learning_rate=0.01)
    t4 = t41 - time.process_time()
    loss_sgd = calculate_squared_loss(X_train, y_train, theta_sgd)
    print("Squared Loss:", loss_sgd)
    print("Theta", np.linalg.norm(theta_sgd))
    print("\n")

    print("Stochastic Gradient Decent with LR=0.05:")
    t51 = time.process_time()
    theta_sgd1 = ls_stochastic_gradient_descent(X_train, y_train, learning_rate=0.05)
    t5 = t51 - time.process_time()
    loss_sgd1 = calculate_squared_loss(X_train, y_train, theta_sgd1)
    print("Squared Loss:", loss_sgd1)
    print("Theta", np.linalg.norm(theta_sgd1))
    print("\n")

    print("Stochastic Gradient Decent with LR=\"adaptive\":")
    t61 = time.process_time()
    theta_sgd2 = ls_stochastic_gradient_descent(X_train, y_train, learning_rate='adaptive')
    t6 = t61 - time.process_time()
    loss_sgd2 = calculate_squared_loss(X_train, y_train, theta_sgd2)
    print("Squared Loss:", loss_sgd2)
    print("Theta", np.linalg.norm(theta_sgd2))
    print("\n")

    print("Done!")
```

```

import numpy as np
from helper import load_data

def sigmoid(z):
    """
    Implements the sigmoid function..
    Args:
        z: A scalar or numpy array of any size
    """
    # handles scalars
    z = np.asarray(z)
    return 1.0 / (1.0 + np.exp(-z))

def logistic_stochastic_gradient_descent(X, y, lr=0.0001):
    """
    Implements the Stochastic Gradient Descent (SGD) algorithm for logistic regression.
    Note:
        - Please do not shuffle your data points.
        - Please use the stopping criteria: number of epochs == 10,000

    Args:
        X: np.array, shape (n, d)
        y: np.array, shape (n,)
        lr: the learning rate for the algorithm

    Returns:
        theta: np.array, shape (d+1,) including the offset term.
    """
    n, d = X.shape
    theta = np.zeros(d + 1)
    step = lr

    Xbar = np.hstack([np.ones((n, 1)), X])

    for _ in range(10000):
        # no shuffle
        for xi, yi in zip(Xbar, y):
            theta_dot_x = np.dot(theta, xi)
            gradient = -1 * yi * xi * sigmoid(-yi * theta_dot_x)
            theta -= step * gradient

    return theta

def stochastic_newtons_method(X, y, lr=0.0001):
    """
    Implements the Stochastic Newton's Method algorithm for logistic regression.
    Note:
        - Please do not shuffle your data points.
        - Please use the stopping criteria: number of epochs == 1,000

    Args:
        X: np.array, shape (n, d)
        y: np.array, shape (n,)
        lr: the learning rate for the algorithm

    Returns:
        theta: np.array, shape (d+1,) including the offset term.
    """
    n, d = X.shape
    theta = np.zeros(d + 1)
    step = lr

    Xbar = np.hstack([np.ones((n, 1)), X])
    eps = 1e-10

    #epochs = 1000
    epochs = 2000

    for _ in range(epochs):
        # no shuffle
        for xi, yi in zip(Xbar, y):
            margin = np.dot(theta, xi)
            sigmoid_signed = sigmoid(yi * margin)

            beta = -1 * yi * (1.0 - sigmoid_signed)
            alpha = sigmoid_signed * (1.0 - sigmoid_signed)

            # add eps to avoid 0 denom
            newton_step = (beta / (alpha * (xi @ xi) + eps)) * xi
            theta -= step * newton_step

    return theta

```